

investigate-a-dataset-template

April 29, 2018

1 Project: Investigate a Dataset (Replace this with something more specific!)

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

Dataset : TMDb movie data This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

Certain columns, like 'cast' and 'genres', contain multiple values separated by pipe (|) characters. There are some odd characters in the 'cast' column. Don't worry about cleaning them. You can leave them as is. The final two columns ending with "_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

Question : Which genres are most popular from year to year? What kinds of properties are associated with movies that have high revenues?

```
In [377]: import numpy as np
import pandas as pd
import seaborn as sn
import matplotlib.pyplot as plt
from pprint import pprint
```

Data Wrangling

1.1.1 General Properties

```
In [378]: # Load data
data = pd.read_csv('data/tmdb-movies.csv')
data.iloc[0]
```

```
Out[378]: id                135397
imdb_id                  tt0369610
popularity                32.9858
budget                 150000000
```

```

revenue                                     1513528810
original_title                             Jurassic World
cast                                       Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
homepage                                 http://www.jurassicworld.com/
director                                 Colin Trevorrow
tagline                                  The park is open.
keywords                                monster|dna|tyrannosaurus rex|velociraptor|island
overview                                Twenty-two years after the events of Jurassic ...
runtime                                  124
genres                                  Action|Adventure|Science Fiction|Thriller
production_companies                    Universal Studios|Amblin Entertainment|Legenda...
release_date                             6/9/15
vote_count                              5562
vote_average                            6.5
release_year                             2015
budget_adj                              1.38e+08
revenue_adj                             1.39245e+09
Name: 0, dtype: object

```

1.1.2 Data Cleaning

Removing duplicate rows if any

```

In [379]: # drop duplicates
pprint("rows before removing duplicates.." + str(len(data)))
data.drop_duplicates(keep = 'first', inplace = True)
pprint("rows after removing duplicates.." + str(len(data)))

#df_new = data.dropna(axis=1, how='all')
#len(df_new)

```

```

'rows before removing duplicates..10866'
'rows after removing duplicates..10865'

```

*** Removing rows who genres are empty***

```

In [380]: dataWithGenres = data[~data['genres'].isnull()]
pprint("rows with genres.." + str(len(dataWithGenres)))
dataWithGenres.head()

```

```

'rows with genres..10842'

```

```

Out[380]:
   id  imdb_id  popularity  budget  revenue \
0  135397  tt0369610   32.985763  150000000  1513528810
1   76341  tt1392190   28.419936  150000000   378436354
2  262500  tt2908446   13.112507  110000000   295238201
3  140607  tt2488496   11.173104  200000000   2068178225

```

4 168259 tt2820852 9.335014 190000000 1506249360

original_title \
0 Jurassic World
1 Mad Max: Fury Road
2 Insurgent
3 Star Wars: The Force Awakens
4 Furious 7

cast \
0 Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1 Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
2 Shailene Woodley|Theo James|Kate Winslet|Ansel...
3 Harrison Ford|Mark Hamill|Carrie Fisher|Adam D...
4 Vin Diesel|Paul Walker|Jason Statham|Michelle ...

homepage director \
0 http://www.jurassicworld.com/ Colin Trevorrow
1 http://www.madmaxmovie.com/ George Miller
2 http://www.thedivergentseries.movie/#insurgent Robert Schwentke
3 http://www.starwars.com/films/star-wars-episod... J.J. Abrams
4 http://www.furious7.com/ James Wan

tagline ... \
0 The park is open. ...
1 What a Lovely Day. ...
2 One Choice Can Destroy You ...
3 Every generation has a story. ...
4 Vengeance Hits Home ...

overview runtime \
0 Twenty-two years after the events of Jurassic ... 124
1 An apocalyptic story set in the furthest reach... 120
2 Beatrice Prior must confront her inner demons ... 119
3 Thirty years after defeating the Galactic Empi... 136
4 Deckard Shaw seeks revenge against Dominic Tor... 137

genres \
0 Action|Adventure|Science Fiction|Thriller
1 Action|Adventure|Science Fiction|Thriller
2 Adventure|Science Fiction|Thriller
3 Action|Adventure|Science Fiction|Fantasy
4 Action|Crime|Thriller

production_companies release_date vote_count \
0 Universal Studios|Amblin Entertainment|Legenda... 6/9/15 5562
1 Village Roadshow Pictures|Kennedy Miller Produ... 5/13/15 6185
2 Summit Entertainment|Mandeville Films|Red Wago... 3/18/15 2480

3	Lucasfilm Truenorth Productions Bad Robot	12/15/15	5292
4	Universal Pictures Original Film Media Rights ...	4/1/15	2947

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

*** Removing rows whose budget and revenue are zero***

```
In [381]: # replace zero to na
dataWithGenres[['budget', 'revenue']] = dataWithGenres[['budget', 'revenue']].replace(
#dataWithGenres.head()
# removing NAN
dataWithGenres.dropna(subset = ['budget', 'revenue'], inplace = True)
rows, col = dataWithGenres.shape
pprint("rows without zero revenue and budget.." + str(len(dataWithGenres)))
```

/opt/conda/lib/python3.6/site-packages/pandas/core/frame.py:2352: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>
self[k1] = value[k2]

'rows without zero revenue and budget..3854'

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>
"""

In [382]: dataWithGenres.dtypes

```
Out[382]: id                int64
imdb_id                  object
popularity              float64
budget                  float64
revenue                 float64
original_title          object
cast                    object
```

```

homepage          object
director          object
tagline           object
keywords          object
overview          object
runtime           int64
genres            object
production_companies object
release_date      object
vote_count        int64
vote_average      float64
release_year      int64
budget_adj        float64
revenue_adj       float64
dtype: object

```

Exploratory Data Analysis

1.1.3 Which genres are most popular?

```
In [383]: genr_arr = []
```


```

genres_data = dataWithGenres['genres'].str.cat(sep = '|')
genres_data_split = pd.Series(genres_data.split('|'))

groupedGenres = genres_data_split.value_counts(ascending = False)
groupedGenres.plot.bar()
plt.xlabel('genres')
#On y-axis
plt.ylabel('Movies Count')
#Name of the graph
plt.title('genres vs movie count')

```

```
Out[383]: Text(0.5,1,'genres vs movie count')
```



Inference : As per the data calculated, movies with ‘drama’ genres occupied the market. Drama is followed by comedy and thriller

1.1.4 Which average budget of movies

```
In [384]: dataWithGenres['budget'].mean()
```

```
Out[384]: 37203696.954852104
```

Inference. The everage budget for a movie is 37203696.954852104

1.1.5 Which movies gives high revenue

```
In [385]: maxRevMovie = dataWithGenres.iloc[dataWithGenres['revenue'].argmax()]
maxRevMovie
```

```
Out[385]: id                9573
imdb_id                tt0309377
popularity                0.545907
budget                5e+07
```

```

revenue                2.61995e+07
original_title          Blood Work
cast                   Clint Eastwood|Jeff Daniels|Anjelica Huston|Wa...
homepage                NaN
director               Clint Eastwood
tagline                He's a heartbeat away from catching the killer
keywords               houseboat|heart|investigation|police|ex-cop
overview               Still recovering from a heart transplant, a re...
runtime                110
genres                 Crime|Drama|Mystery|Thriller
production_companies   Malpaso Productions|Warner Bros. Pictures
release_date           8/4/02
vote_count              89
vote_average           5.7
release_year           2002
budget_adj             6.06131e+07
revenue_adj            3.17607e+07
Name: 4021, dtype: object

```

```
In [386]: ### Average runtime of movies
```


```
In [387]: dataWithGenres['runtime'].mean()
```

```
Out[387]: 109.22029060716139
```

```
In [388]: ### Movies run time plot
```

```
In [389]: data_runtime = dataWithGenres['runtime']
plt.hist(data_runtime, bins=50)
plt.xlabel('Movie Runtime')
#On y-axis
plt.ylabel('Movies Count')
#Name of the graph
plt.title('Runtime of the movies')
```

```
Out[389]: Text(0.5,1,'Runtime of the movies')
```



Inference. The histogram follows positively skewed one where most of the movie runtime falls around 75 to 135.

Will print the statistics based on runtime

```
In [390]: dataWithGenres['runtime'].describe()
```

```
Out[390]: count    3854.000000
          mean      109.220291
          std       19.922820
          min       15.000000
          25%       95.000000
          50%      106.000000
          75%      119.000000
          max       338.000000
          Name: runtime, dtype: float64
```

Inference. Based on runtime statistics, average runtime is 109 and it has standard deviation of 19


```
In [391]: ### Release year Vs Revenue
```

```
In [392]: dataByYear = dataWithGenres.groupby('release_year')['revenue'].sum()
          dataByYear.plot.line()
          plt.xlabel('Year of release')
          #On y-axis
```



```
plt.ylabel('Revenue')
#Name of the graph
plt.title('Release Year vs revenue')
```

Out[392]: Text(0.5,1,'Release Year vs revenue')



Inference : The maximum revenue is on 2015

```
In [393]: dataByYear.describe()
```

```
Out[393]: count      5.600000e+01
          mean       7.411147e+09
          std        8.242757e+09
          min        8.473669e+07
          25%        7.454036e+08
          50%        2.942177e+09
          75%        1.353885e+10
          max        2.620292e+10
          Name: revenue, dtype: float64
```

###Runtime Vs revenue


```
In [394]: datarunRev = dataWithGenres.groupby('runtime')['revenue'].sum()
          datarunRev.plot.line()
          plt.xlabel('run time of movie')
```

```

#On y-axis
plt.ylabel('Revenue')
#Name of the graph
plt.title('runtime vs revenue')

```

Out[394]: Text(0.5,1,'runtime vs revenue')



Inference : Based on the graph above, the graph is positively skewed where the revenue is contained by the runtime time. After runtime of 130.. the revenue is decreasing.

1.1.6 What kinds of properties are associated with movies that have high revenues?

In [395]: *#To find the paroperties that are associated with high revenues. We have to find the c*
`data.corr()`

Out[395]:

	id	popularity	budget	revenue	runtime	vote_count	\
id	1.000000	-0.014351	-0.141341	-0.099235	-0.088368	-0.035555	
popularity	-0.014351	1.000000	0.545481	0.663360	0.139032	0.800828	
budget	-0.141341	0.545481	1.000000	0.734928	0.191300	0.632719	
revenue	-0.099235	0.663360	0.734928	1.000000	0.162830	0.791174	
runtime	-0.088368	0.139032	0.191300	0.162830	1.000000	0.163273	
vote_count	-0.035555	0.800828	0.632719	0.791174	0.163273	1.000000	
vote_average	-0.058391	0.209517	0.081067	0.172541	0.156813	0.253818	
release_year	0.511393	0.089806	0.115904	0.057070	-0.117187	0.107962	
budget_adj	-0.189008	0.513555	0.968963	0.706446	0.221127	0.587062	

revenue_adj	-0.138487	0.609085	0.622531	0.919109	0.175668	0.707941
-------------	-----------	----------	----------	----------	----------	----------

	vote_average	release_year	budget_adj	revenue_adj
id	-0.058391	0.511393	-0.189008	-0.138487
popularity	0.209517	0.089806	0.513555	0.609085
budget	0.081067	0.115904	0.968963	0.622531
revenue	0.172541	0.057070	0.706446	0.919109
runtime	0.156813	-0.117187	0.221127	0.175668
vote_count	0.253818	0.107962	0.587062	0.707941
vote_average	1.000000	-0.117576	0.093079	0.193062
release_year	-0.117576	1.000000	0.016771	-0.066236
budget_adj	0.093079	0.016771	1.000000	0.646627
revenue_adj	0.193062	-0.066236	0.646627	1.000000

As per the table above, revenue is more correlated with vote_count with the correlation value of 0.79

Inference. Higher the value between columns , higher the correlation is.

```
In [396]: # Compute the correlation matrix
corr = dataWithGenres.corr()


mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

f, ax = plt.subplots(figsize=(11, 9))

cmap = sn.diverging_palette(220, 10, as_cmap=True)

sn.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
           square=True, linewidths=.5, cbar_kws={"shrink": .5})

Out[396]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb50511e908>
```



Conclusions

As per the data of movies, High revenue is associated with the following in the order it was associated. 1. Vote_count (having correlation value of 79%). If the vote_count is high then obvious more people come and see the movie which in turns increase the revenue) 2. Budget (having correlation value of 73%. Higher the budget is , higher the revenue of movie as per the dataset) 3. Popularity (having correlation value of 66.3%. Higher the popularity of movie is , higher the revenue of movie as per the dataset)

As per the analysis following are important to be considered 1. To get more revenue, the run time of movies should be around 109 2. The mean budget of all movies is 37203696. This is average amount to make a good revenue movie 2. As the years moves on, the revenue of movie increases based on the causation mention just above like vote_count, budget, popularity

Limitations : 1. The runtime measure is not shown whether it is second or minutes 2. The measure of budget/revenue are not shown. 3. Few invalid data/duplicate data's has been excluded from analysis. Not sure whether that affect our analysis. We need to get that data corrected.